CV

Greg Nwosu

September 15, 2017

Summary

I am a technology enthusiast with a lot of experience in functional programming and big data technologies. I've some commercial experience in machine learning but the majority of my experience is from personal projects and online courses. I've just arrived back after travelling for 6 months in SE Asia touring places of cultural and technological interest with my family. I am certified in machine learning from Stanford University's course. I'm currently half way through Udacity's Deep Learning Foundation Course with a view to complete the Udacity's Self Driving Car NanoDegree next year. My goal is to achieve the rare position of being adept in machine learning and data science together with my current skills of big data engineering and ETL.

Current Interests

My current interests are almost exclusively in FP, including:

- Embedded DSL's using the Fixed Point Combinator in Haskell,
- Purescript S.P.A.'s
- Idris
- Bayesian Methods for Data Analytics
- Machine Learning , LSTM Generative Nets
- Parsec like Parsing
- General chore automation, planning on creating Internet of things devices.
- Haskell FFI - recently created and open sourced the blink1device Haskell FFI API

Other hobbies include:

- Cycling
- studying civil litigation and human rights law
- Zen Meditation.
- I have also trained to teach adults and am a mentor to children learning to code at CoderDojo.

Commercial & Technical Experience

July 2016 - April 2017, Self Employed: Aimia

Machine Learning Engineer

Responsibilities

Working within a new team for monetising Aimia's vast data repository. My responsibilities initially were helping migrate a Aimia service to a more robust framework (EMR). (This wasn't within my remit; but as the new team was not yet properly formed , I rolled up my sleeves and helped out.) I then moved into co-designing and developing a platform to capture all the data needed for the machine learning techniques we wished to use. This completed, I then developed real-time ETL spark streams for data out of the legacy hive data-warehouse so data could easily be used in a variety of machine learning algorithims. Choosing parquet because of its schema evolution and performance properties. At this stage I was offered the job of lead engineer on the team which I turned down because my wife and children wanted to go travelling. I completed

Commercial & Technical Experience Summary

skill	years
FIX	2
Trading	2
Leadership, Mentoring	4
Financial Reconciliation	2
Sports Betting	1
M-Commerce	1
Dev-ops	1
Consultancy	3
Government	3
Financial Markets	5
Securities, FX	3
Scala	3
Big Data	3
Machine Learning	1
Spark	2
Haskell	2
Spray	0.5

Education

2002-2003 University College London

M.Sc. Intelligent Systems (Incomplete)

Course Content

- Neural Networks
- SVMs
- Decision Trees
- Learning theory
- Maximum Likelihood Estimation
- Bayesian Decision Theory
- Hidden Markov Models
- EM Algorithm
- ICA
- Clustering
- Factor Analysis